Machine Learning End Term Exam

Ahmad Shayaan IMT2014004
H Vijaya Sharvani IMT2014022
Indu Ilanchezian IMT2014024

December 15, 2017

Question 2

To derive the solution to the modified linear regression and to show that it leads to the generalized form of ridge regression.

Solution:-

Given the attribute $x_i = \hat{x_i} + \epsilon_i$, where the $\hat{x_i}$ is the true measurement and ϵ_i is the zero mean vector with covariance matrix $\sigma^2 I$.
Modified loss function

$$W^* = argmin_W E_\epsilon \sum_{i=1}^{n} (y_i - W^T(\hat{x_i} + \epsilon_i))^2$$

Where W is the transformation vector.

$$W^* = argmin_W E_\epsilon ||Y - (X + \epsilon)W||_2^2 \tag{1}$$

Where

$$Y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$X = \begin{bmatrix} \hat{x}_1^T \\ \hat{x}_2^T \\ \vdots \\ \hat{x}_n^T \end{bmatrix}$$

$$\epsilon = \begin{bmatrix} \epsilon_1^T \\ \epsilon_2^T \\ \vdots \\ \epsilon_n^T \end{bmatrix}$$

Expanding right hand side of equation (1).

$$E_\epsilon||Y-(X+\epsilon)W||_2^2 = E_\epsilon\Big[(Y-(X+\epsilon)W)^T(Y-(X+\epsilon)W)\Big]$$

$$= E_\epsilon\Big[Y^TY + W^T(X+\epsilon)^T(X+E) - 2W^T(X+E)^TY\Big] \tag{2}$$

To minimize the equation we will differentiate equation (2) with respect to W.

$$\frac{\partial E_\epsilon\Big[Y^TY + W^T(X+\epsilon)^T(X+\epsilon)W - 2W^T(X+E)^TY\Big]}{\partial W} = 0$$

We know that $\frac{\partial E(f(x))}{\partial} = E\frac{\partial f(x)}{\partial x}$.

$$E_\epsilon\Big[\frac{\partial Y^TY}{\partial W} + \frac{\partial W^T(X+\epsilon)^T(X+\epsilon)W}{\partial W} - 2\frac{\partial W^T(X+E)^TY}{\partial W}\Big] = 0$$

$$E_\epsilon\big[2(X+\epsilon)^T(X+\epsilon)W - 2(X+\epsilon)^TY\big] = 0$$

$$2E_\epsilon\big[(X+\epsilon)^T(X+\epsilon)W\big] - 2E_\epsilon\big[(X+\epsilon)^TY\big] = 0$$

$$E_\epsilon\big[(X^TX + \epsilon^T\epsilon + 2\epsilon^TX)W\big] = E_\epsilon\big[(X+\epsilon)^TY\big])$$

$$E_\epsilon(X^TXW) + E_\epsilon(\epsilon^T\epsilon W) + 2E_\epsilon(\epsilon^TXW) = E_\epsilon(X^TY) + E_\epsilon(\epsilon^TY)$$

We know that E(AB) = E(A)E(B) if A and B are independent variables and $E_f(h(x)) = \int_{-\infty}^{\infty} h(x)f(x)dx$.

$$\sum_{i=1}^{n} X^TXWP(\epsilon_i) + E_\epsilon(\epsilon\epsilon^T)E_\epsilon(W) + 2E_\epsilon(X)E_\epsilon(\epsilon) = \sum_{i=1}^{n} X^TYP(\epsilon_i) + E_\epsilon(Y)E_\epsilon(\epsilon)$$

We know that the noise is a zero mean Gaussian noise therefore $E_\epsilon(\epsilon) = 0$

$$(X^TX + \sigma^2I)W = X^TY$$

$$W = (X^TX + \sigma^2I)^{-1}X^TY$$

Therefore the solution of the minimization is

$$W^* = (X^TX + \sigma^2I)^{-1}X^TY$$

This solution is same as the solution for Ridge regression.

$$W^* = (X^TX + \lambda I)^{-1}X^TY$$

Question 3

$VC(\mathcal{H})$ is the maximum cardinality of any set of instances that can be shattered by $\mathcal{H}$. We say that $\mathcal{H}$ shatters a set of points if and only if it can assign any possible labeling to those points.

1. We should show that the VC dimension $d_{\mathcal{H}}$ of any finite hypothesis space $\mathcal{H}$ is at most $log_2\mathcal{H}$.
   Proof:
   For any set of distinct points S of size m, there are 2^m distinct ways of labeling those points. This means that for $\mathcal{H}$ to shatter S it must contain at least 2^m distinct hypotheses. This tells us that if the VC dimension of $\mathcal{H}$ is m then we must have 2^m hypotheses, i.e. $2^m \leq |\mathcal{H}|$ or equivalently that $m = VC(\mathcal{H}) \leq log_2|\mathcal{H}|$.

2. We should show that the bound for VC dimension is tight and for any $d > 1$, there exists a hypothesis class $\mathcal{H}$ such that $d = d_{\mathcal{H}}$.
   Proof:
   Consider a domain with n binary features and binary class labels. Let $\mathcal{H}$ be the hypothesis space that contains all decision trees over those features that have depth no greater than d_e. (The depth of a decision tree is the depth of the deepest leaf node.)
   First note that any tree in $\mathcal{H}$ can be represented by a tree of exactly depth d_e in $\mathcal{H}$. So we will restrict our attention to trees of exactly depth d_e. All of these trees have 2^{d_e} leaf nodes. Also note that there are a total of 2^n examples in our instance space, which gives us an immediate upper bound on the VC-dimension of $\mathcal{H}$, i.e. $VC(\mathcal{H}) \leq 2^n$.
   To get a lower bound let S contain the set of all possible 2^n instances. Since we have that $d_e \geq n$ it is straightforward to create a tree of depth n with a leaf node for each example and furthermore we can label the leaf nodes in all possible ways. This shows that we can shatter the set S with $\mathcal{H}$, which implies that $VC(\mathcal{H}) \geq 2^n$. Combining the upper and lower bound tell us that $VC(\mathcal{H}) = 2^n$, i.e. $d_{\mathcal{H}} = 2^n$. Hence, we have showed a tight bound on the VC dimension of hypothesis space $\mathcal{H}$.

   Now we will show that for any $d > 1$, there exists a hypothesis class $\mathcal{H}$ such that $d = d_{\mathcal{H}}$.
   Take a set of size d, $C = \{e_1, e_2, ..., e_d\}$ such that $\{e_i; i \in [d]\}$ is the standard basis in R^d. To prove that C shatters $\mathcal{H}$, it suffices to show that $|\mathcal{H}_C| = 2^d$. The hypothesis on the set C is given as,

$$\mathcal{H}_C = \{h(c), h \in HS_d\} = \{h(e_1, h(e_2), ..., h(e_d), h \in HS_d\}$$

For a particular $\omega^T = (\omega_1, \omega_2, ..., \omega_d)$,

$$\begin{aligned} h(c) &= (\langle \omega, c \rangle, c \in C) \\ &= (\langle \omega_1, e_1 \rangle, \langle \omega_2, e_2 \rangle, ..., \langle \omega_d, e_d \rangle) \\ &= (\omega_1, \omega_2, ..., \omega_d) \\ &= \omega \end{aligned}$$

Since all possible combinations 2^d be chosen on ω.

$$\begin{aligned} &\implies |\mathcal{H}_C| = 2^d \\ &\implies VCdim(HS_d) \geq d \end{aligned}$$

Let us take a arbitrary set C of size d + 1.

$$C = \{x_1, x_2, ..., x_{d+1}\}, x_i \in R^d$$

Since, $x_i's$ are coming from d dimensional space, $\{x_i, i \in [d+1]\}$ are linearly dependent,

$$\implies \exists\ a_1, a_2, ..., a_{d+1}\ s.t.\ \sum_{i=1}^{d+1} a_i x_i = 0$$

Let $I = \{i, a_i > 0\}\ and\ J = j, a_j > 0$

$$\implies \sum_{i \in I} a_i x_i = -\sum_{j \in J} a_j x_j = \sum_{j \in J} |a_j| x_j$$

Suppose C is shattered by $\mathcal{H}$.

$$Claim : \exists\ \omega\ s.t\ \langle \omega, x_i \rangle > 0\ \forall i \in I\ \&\ \langle \omega, x_j \rangle < 0, \forall j \in J$$

$$\begin{aligned} \implies 0 &< \sum_{i \in I} a_i \langle \omega, x_i \rangle \\ &= \sum_{i \in I} \langle \omega, a_i x_i \rangle \\ &= \sum_{j \in J} \langle \omega, |a_j| x_j \rangle \\ &= \sum_{j \in J} |a_j| \langle \omega, x_j \rangle < 0, which\ is\ a\ contradiction \end{aligned}$$

Therefore $|\mathcal{H}_C| < 2^{d+1}$ for any arbitrary set of size d + 1. So, the VC-dimension of the class of homogeneous halfspace in R^d is d.

Question 4

Algorithm 1 Building Decision tree

procedure BUILDDECISIONTREE(Dataset D,Target_Attributes,Attributes)
    T = ϕ ▷ Initializing a tree
    if All Target attributes of one type **then**
        return a node with single label
    else if Attributes = ϕ **then**
        return single node tree i.e. root with label = most common value of the target attribute in the dataset
    else
        Attributes* = GETCHISQUARESCORE(Attributes)
        A = GETINFOGAIN(Attributes*,Target_Attributes)
        for all possible values of A **do**
            if A has missing values **then**
                n = Total number of data points in A
                n^* = Number of non-missing values
                $\mu_i = \frac{n_1^*}{n^*} \dots \frac{n_{m_i}^*}{n*}$
                $n_i = n_i^* + \mu_i(n - n^*)$
                D[A].append(n_i) ▷ Adding missing values of attribute A
                missing_probab.append(μ_i) ▷ Store the missing value probabilities
            subset = The set of data points with value v_i for A
            T.addNode[BuildDecisionTree(subset,Target_Attributes,Attributes-A)]
    return T
function GETCHISQUARESCORE(Attributes)
    l = [ ] ▷ Empty list initialization
    for A in Attributes **do**
        $\mu_i = \frac{n_i}{n}$ ▷ n_i: number of data points with i^{th} label
        n_{ij} = number of points with label i in partition j
        $e_{ij} = \mu_i \sum_i n_{ij}$
        score = $\sum_i \frac{(n_{ij}-e_{ij})^2}{e_{ij}}$
        if score $\geq$ threshold **then** ▷ Correlation for partition j
            l.append(A)
    return l ▷ list containing attributes with score greater than threshold
function GETINFOGAIN(Attributes*,Target_Attributes)
    Target_Entropy = -$\sum_i P(y=i)\log(P(y=i))$ ▷ Entropy of the target variable
    for all i in Attributes* **do**
        ▷ $P(y=k|x_i=v_{ij}) = \frac{\theta_{ijk}P(y=k)}{P(x_i=v_{ij})}$
        Attributes_Entropy = -$\sum_{ijk} \frac{\theta_{ijk}P(y=k)}{P(x_i=V_{ij})} \log\left[\frac{\theta_{ijK}P(y=k)}{P(x_i=V_{ij})}\right]$
        gain = 0
        gain = Target_Entropy - Attribute_Entropy
    return Attribute with max information Gain

Algorithm 2 Testing Decision Tree

```
1: procedure testTree(Tree T, Datpoint D, missing_probab)
2:                                   ▷ Node is a list that stores all the leaf nodes
3:     if T.node == Leaf or D == φ then
4:         return T.node
5:     for all attributes A in D do                          ▷ A is an attribute of D
6:         if A is not in T.node then
7:             continue
8:         else
9:             if value(A) != φ then
10:                d = T.checknode(value(A))
11:                                        ▷ d is the decision taken at that node
12:                Node.append[testTree(T.takepath(d), D-A, missing_probab)]
13:                        ▷ Take path along branch taken according to decision d
14:            else if value(A) == φ then
15:                for All probabilities p in missing_probab do
16:                        ▷ probabilities stored for missing values of attributes at the time of testing
17:                    Push decision tree in each branch of the node with probability μ_i
18:                    Node.append(all the returned leaf nodes)
19:        if Node.size == 1 then
20:            return Probability 1 for that class and zero for other
21:        else
22:            S = Sum of probabilities for each class
23:                    ▷ S is a list that contains class membership probabilities
24:            return S
```

Question 5

Let $\vec{x_1}, \vec{x_2}...\vec{x_n}$ be the feature vectors of n data points in the original feature space. Let ϕ be the feature transformation function. Then, $\phi(\vec{x_1}), \phi(\vec{x_2})..., \phi(\vec{x_n})$ are the feature vectors in the transformed feature space.
Let K be the kernel function such that:

$$K(i,j) = \phi(x_i)^T \phi(x_j)$$

The center of mass, $\vec{\mu}$, in the feature space can be defined as the average of the vectors in the transformed feature space.

$$\vec{\mu} = \frac{1}{n} \sum_{i=1}^{n} \phi(\vec{x_i})$$

Consider:

$$\begin{aligned}
||\mu||^2 &= \mu^T \mu \\
&= \mu^T \frac{1}{n} \sum_{i=1}^{n} \phi(\vec{x_i}) \\
&= \frac{1}{n} \sum_{j=1}^{n} \phi(\vec{x_j})^T \frac{1}{n} \sum_{i=1}^{n} \phi(\vec{x_i}) \\
&= \frac{1}{n^2} \sum_{i,j} \phi(\vec{x_j})^T \phi(\vec{x_i}) \\
&= \frac{1}{n^2} \sum_{i,j} K(i,j)
\end{aligned}$$

Average of the squared Euclidean distances from μ to each $\phi(x)$

The squared euclidean distance of a single feature vector in the transformed space from the center of mass $\vec{\mu}$ can be expressed as follows:

$$\begin{aligned}
||\phi(\vec{x_i}) - \vec{\mu}||^2 &= (\phi(\vec{x_i}) - \vec{\mu})^T (\phi(\vec{x_i}) - \vec{\mu}) \\
&= \phi(\vec{x_i})^T \phi(\vec{x_i}) - 2\phi(\vec{x_i})^T \vec{\mu} + ||\vec{\mu}||^2 \\
&= K(i,i) - \frac{2}{n} \phi(\vec{x_i})^T \sum_{j=1}^{n} \phi(\vec{x_j}) + ||\vec{\mu}||^2 \\
&= K(i,i) - \frac{2}{n} \sum_{j=1}^{n} \phi(\vec{x_i})^T \phi(\vec{x_j}) + ||\vec{\mu}||^2 \\
&= K(i,i) - \frac{2}{n} \sum_{j=1}^{n} K(i,j) + \frac{1}{n^2} \sum_{r,s} K(r,s)
\end{aligned}$$

The average of the euclidean distances of all the points from the center of mass can be written as:

$$\begin{aligned}
\frac{1}{n} \sum_{i=1}^{n} ||\phi(\vec{x_i}) - \vec{\mu}||^2 &= \frac{1}{n} \left(\sum_{i=1}^{n} \left(K(i,i) - \frac{2}{n} \sum_{j=1}^{n} K(i,j) + \frac{1}{n^2} \sum_{r,s} K(r,s) \right) \right) \\
&= \frac{1}{n} \left(\sum_{i=1}^{n} K(i,i) - \frac{2}{n} \sum_{i,j} K(i,j) + \frac{n}{n^2} \sum_{r,s} K(r,s) \right) \\
&= \frac{1}{n} \left(\sum_{i=1}^{n} K(i,i) - \frac{1}{n} \sum_{i,j} K(i,j) \right)
\end{aligned}$$

Thus, the average of euclidean distances from the center of mass $\vec{\mu}$ to each $\phi(x)$ can be expressed in terms of the kernel function K.